


Future-State Design

Meridian Group Data Platform Consolidation

Target data model

One relational schema (see schema.sql) with a single canonical vendor record per real-world vendor, a crosswalk table that keeps each entity's original legacy ID for audit purposes, and unified invoice and bank transaction tables keyed to canonical vendor and entity IDs. Nothing about the legacy IDs gets thrown away — they just stop being the primary way the data is organized.

Reconciliation, before and after



The four workstreams

Workstream	Deliverable
Vendor master consolidation	migrate_and_cleanse.py — normalizes vendor names, resolves duplicates, keeps the legacy-ID crosswalk
Vendor setup validation	vendor_setup_validation.py — blocks a vendor from the payment cycle until tax ID, terms, and status all check out
Exception-based reconciliation	reconciliation_queries.sql + invoice_integration_simulation.py — surfaces only unmatched items and variances
KPI reporting	the dashboard, plus a documented Power BI equivalent

Design principles

One canonical record per real-world vendor or invoice, not three versions of the same thing. Validation happens at the point of entry, not the point of failure — a vendor with no tax ID should never make it as far as the payment queue. Reconciliation is exception-based rather than exhaustive; a process that makes someone look at every line doesn't scale, and one that only surfaces what's actually wrong does. Every migrated row is traceable back to its legacy source through the crosswalk table, so de-duplicating never means losing a fact.

Interfaces

The dashboard reads a single JSON export produced by the SQL/Python pipeline, so there's no manual copy-paste between systems. In a production deployment, Power BI would connect to the same SQL layer directly — the Power Query and DAX for that path are documented separately.