

Data Migration Plan

Meridian Group Data Platform Consolidation

Approach

Extract, cleanse, resolve, load — run as a single, repeatable pass over the three entities' source files rather than a series of manual steps someone has to remember to do in order.

Step	What happens
Extract	Read each entity's vendor, invoice, and bank files as they are.
Cleanse	Parse each entity's own date format; strip currency symbols and commas from amounts; normalize payment-terms notation to a single integer.
Resolve	Normalize vendor name casing and whitespace, match against existing canonical vendors, assign one canonical vendor ID, and write the legacy-ID crosswalk.
Flag	Mark exact-duplicate invoice entries (same entity, vendor, amount, and date) rather than silently dropping or silently keeping them.
Load	Write to the target database and mirror the result to clean CSVs for human review.

Entity-specific quirks handled

Entity	Quirk	Handling
Meridian North	Amounts as plain numbers, dates as YYYY-MM-DD	Parsed natively
Meridian West	Amounts as "\$12,508.72", dates as MM/DD/YYYY	Currency symbols and commas stripped before conversion
Meridian Coastal	Dates as DD-Mon-YYYY, active flag as 1/0 instead of Y/N	Entity-specific date parser; flag normalizer accepts Y/N, Yes/No, and 1/0

Validation after load

- 1 Row counts reconciled between legacy and loaded data, accounting for intentional duplicate-flagging rather than silent drops.
- 2 Every crosswalk row resolves to exactly one canonical vendor.
- 3 Every invoice and bank transaction references a valid entity, and every invoice a valid vendor.
- 4 Vendor validation run against the migrated vendor master and reviewed before sign-off.

Rollback

The migration is a deterministic function of the source files — the database gets rebuilt from scratch on every run rather than patched in place. If something's wrong, the fix is to correct the source data or the transform logic and run it again. Nothing in the process touches the legacy files themselves.